

Building a Clean Architecture Banking Application

Spring Boot 3.5 & PostgreSQL

A Production-Ready RESTful Banking API

By: Wildan Anugrah

Github

<https://github.com/wildananugrah/spring-learning>

By: Wildan Anugrah

Table of Contents

Contents

Note: In HTML export, these links are clickable. In PDF, use page navigation.

1. **What We'll Build** - Application overview
2. **Architecture Overview** - Three-layered design
3. **Step 1: Project Setup** - Dependencies & configuration
4. **Step 2: Entity Models** - Database schema
5. **Step 3: DTOs** - Data transfer objects
6. **Step 4: Exception Handling** - Global error management
7. **Step 5: Repository Layer** - Data access with JPA
8. **Step 6: JWT Authentication** - Secure token-based auth
9. **Step 7: Logic Layer** - Business logic implementation
10. **Step 8: Route Layer** - Controllers & API endpoints

Contents (cont.)

11. **Step 9: Docker Deployment** - Containerization
12. **Step 10: Testing the Application** - Running & testing
13. **Step 11: Spring Boot Actuator** - Production monitoring
14. **Step 12: Comprehensive Logging** - Request/response logging with JSON support
15. **Key Features Implemented** - What we've built
16. **Best Practices Demonstrated** - Professional development
17. **Project Structure Recap** - Clean organization
18. **Conclusion** - Summary & next steps
19. **Resources** - Documentation & testing
20. **Bonus: N+1 Query Problem** - Performance optimization guide

 **Tip:** Export to HTML for clickable navigation, or use PDF bookmarks feature

What We'll Build

A Fully Functional Banking Application

Application Features

Core Functionality

- User registration and authentication with JWT
- Account creation and management
- Deposit, withdrawal, and transfer operations
- Transaction history with filtering, sorting, and pagination
- Spring Boot Actuator for application monitoring

Architecture

- Clean separation of concerns with three distinct layers
- RESTful API design
- Secure JWT authentication
- PostgreSQL database with HikariCP connection pooling

Architecture Overview

Clean Three-Layered Architecture

The Three Layers

1. Route Layer (`/routes` folder)

Purpose: Handle HTTP requests and responses

Responsibilities:

- Receive incoming requests
- Validate input using Bean Validation
- Send formatted responses
- Handle authentication context

The Three Layers (cont.)

2. Logic Layer (`/logics` folder)

Purpose: Implement business logic

Responsibilities:

- Process business rules
- Coordinate between routes and services
- Handle transactions and data transformations
- Enforce business constraints

The Three Layers (cont.)

3. Service Layer (`/services` folder)

Purpose: Manage dependencies and data access

Responsibilities:

- Database operations (JPA repositories)
- External service integrations
- Configuration management

Architecture Benefits

Why This Separation?

- ✓ **Maintainability:** Each layer has a single responsibility
- ✓ **Testability:** Layers can be tested independently
- ✓ **Scalability:** Easy to modify or replace individual layers
- ✓ **Readability:** Clear code organization

Project Structure

```
src/main/java/com/user/account/app/
  config/          # Configuration classes
    JwtUtil.java
    JwtAuthenticationFilter.java
    SecurityConfig.java
  dto/            # Data Transfer Objects
  entities/        # JPA Entities
  exceptions/      # Custom Exceptions
  logics/          # Business Logic Layer
  routes/          # Controllers (Route Layer)
  services/        # Data Access Layer
```

Step 1: Project Setup

Getting Started

Prerequisites

-  Java 17 or higher
-  Maven 3.6+
-  PostgreSQL database
-  Your favorite IDE (IntelliJ IDEA, VS Code, etc.)
-  Basic knowledge of Spring Boot and REST APIs

Key Dependencies

```
<!-- Web -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-web</artifactId>
</dependency>

<!-- JPA & Database -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-data-jpa</artifactId>
</dependency>
<dependency>
    <groupId>org.postgresql</groupId>
    <artifactId>postgresql</artifactId>
</dependency>

<!-- Security -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-security</artifactId>
</dependency>
```

More Dependencies

```
<!-- Validation -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-validation</artifactId>
</dependency>

<!-- Actuator for Monitoring -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-actuator</artifactId>
</dependency>

<!-- Logstash for JSON Logging -->
<dependency>
    <groupId>net.logstash.logback</groupId>
    <artifactId>logstash-logback-encoder</artifactId>
    <version>8.0</version>
</dependency>

<!-- Lombok -->
<dependency>
    <groupId>org.projectlombok</groupId>
    <artifactId>lombok</artifactId>
    <optional>true</optional>
</dependency>
```

JWT Dependencies

```
<!-- JWT Authentication -->
<dependency>
    <groupId>io.jsonwebtoken</groupId>
    <artifactId>jjwt-api</artifactId>
    <version>0.12.3</version>
</dependency>
<dependency>
    <groupId>io.jsonwebtoken</groupId>
    <artifactId>jjwt-impl</artifactId>
    <version>0.12.3</version>
    <scope>runtime</scope>
</dependency>
<dependency>
    <groupId>io.jsonwebtoken</groupId>
    <artifactId>jjwt-jackson</artifactId>
    <version>0.12.3</version>
    <scope>runtime</scope>
</dependency>
```

Maven Compiler Plugin

```
<plugin>
  <groupId>org.apache.maven.plugins</groupId>
  <artifactId>maven-compiler-plugin</artifactId>
  <version>3.14.1</version>
  <configuration>
    <source>17</source>
    <target>17</target>
    <annotationProcessorPaths>
      <path>
        <groupId>org.projectlombok</groupId>
        <artifactId>lombok</artifactId>
        <version>1.18.36</version>
      </path>
    </annotationProcessorPaths>
  </configuration>
</plugin>
```

Important: This ensures Lombok annotations are processed correctly!

Database Configuration

```
# Application Info
spring.application.name=${APP_NAME:user-account}
server.port=${APP_PORT:8000}

# Database Configuration (using environment variables)
spring.datasource.url=${SPRING_DATASOURCE_URL:jdbc:postgresql://localhost:6435/bank_app}
spring.datasource.username=${SPRING_DATASOURCE_USERNAME:pg}
spring.datasource.password=${SPRING_DATASOURCE_PASSWORD:p@ssw0rd1234}
spring.datasource.driver-class-name=org.postgresql.Driver

# HikariCP Connection Pool
spring.datasource.hikari.minimum-idle=5
spring.datasource.hikari.maximum-pool-size=20
spring.datasource.hikari.idle-timeout=300000
spring.datasource.hikari.max-lifetime=1800000
spring.datasource.hikari.connection-timeout=30000
```

Note:

- Application runs on port **8000** (not 8080)
- Uses environment variables with fallback defaults

JPA Configuration

```
# JPA Configuration
spring.jpa.hibernate.ddl-auto=update
spring.jpa.show-sql=true
spring.jpa.properties.hibernate.format_sql=true
spring.jpa.properties.hibernate.dialect=org.hibernate.dialect.PostgreSQLDialect

# JWT Configuration
jwt.secret=YOUR_256_BYTE_SECRET_HERE
jwt.expiration=86400000

# Logging
logging.level.org.springframework.security=DEBUG
```

Generate JWT Secret: `openssl rand -hex 256`

Actuator Configuration

```
# Actuator Configuration
management.endpoints.web.exposure.include=health,info,metrics,env,beans,mappings
management.endpoint.health.show-details=when-authorized
management.health.db.enabled=true
management.info.env.enabled=true

# Application Info
info.app.name=Bank Application API
info.app.description=User Account Management System
info.app.version=1.0.0
```

Step 2: Entity Models

Database Schema Design

User Entity

```
@Entity
@Table(name = "users")
@Data
@Builder
@NoArgsConstructor
@AllArgsConstructor
public class User {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    @Column(nullable = false, unique = true)
    private String username;

    @Column(nullable = false, unique = true)
    private String email;

    @Column(nullable = false)
    private String password;

    @Column(nullable = false)
    private String fullName;
```

User Entity (cont.)

```
@OneToMany(mappedBy = "user", cascade = CascadeType.ALL)
private List<Account> accounts;

@CreationTimestamp
@Column(nullable = false, updatable = false)
private LocalDateTime createdAt;

@UpdateTimestamp
@Column(nullable = false)
private LocalDateTime updatedAt;
}
```

Key Features:

- Lombok `@Data`, `@Builder` for clean code
- Automatic timestamp management
- One-to-many relationship with accounts

Account Entity

```
@Entity
@Table(name = "accounts")
@Data
@Builder
@NoArgsConstructor
@AllArgsConstructor
public class Account {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    @Column(nullable = false, unique = true)
    private String accountNumber;

    @Column(nullable = false)
    private String accountName;

    @Column(nullable = false, precision = 19, scale = 2)
    private BigDecimal balance;
```

Account Entity (cont.)

```
@ManyToOne(fetch = FetchType.LAZY)
@JoinColumn(name = "user_id", nullable = false)
private User user;

@OneToMany(mappedBy = "account", cascade = CascadeType.ALL)
private List<Transaction> transactions;

@CreationTimestamp
@Column(nullable = false, updatable = false)
private LocalDateTime createdAt;

@UpdateTimestamp
@Column(nullable = false)
private LocalDateTime updatedAt;
}
```

Key Features:

- `BigDecimal` for precise currency calculations
- Lazy loading for performance
- Bidirectional relationships

Transaction Entity

```
@Entity
@Table(name = "transactions")
@Data
@Builder
@NoArgsConstructor
@AllArgsConstructor
public class Transaction {
    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    private Long id;

    @ManyToOne(fetch = FetchType.LAZY)
    @JoinColumn(name = "account_id", nullable = false)
    private Account account;

    @Enumerated(EnumType.STRING)
    @Column(nullable = false)
    private TransactionType type;

    @Column(nullable = false, precision = 19, scale = 2)
    private BigDecimal amount;
```

Transaction Entity (cont.)

```
@Column(precision = 19, scale = 2)
private BigDecimal balanceBefore;

@Column(precision = 19, scale = 2)
private BigDecimal balanceAfter;

@Column
private String description;

@Column
private String referenceNumber;

@Column
private String toAccountNumber;

@Column
private String fromAccountNumber;

@CreationTimestamp
@Column(nullable = false, updatable = false)
private LocalDateTime createdAt;
```

Transaction Type Enum

```
public enum TransactionType {  
    DEPOSIT,  
    WITHDRAWAL,  
    TRANSFER_IN,  
    TRANSFER_OUT  
}
```

Key Features:

- Tracks balance before/after each transaction
- Reference numbers for audit trail
- Transfer tracking with account numbers
- Enum for type safety

Step 3: DTOs

Data Transfer Objects

Authentication DTOs

RegisterRequest

```
@Data  
@Builder  
@NoArgsConstructor  
@AllArgsConstructor  
public class RegisterRequest {  
    @NotBlank(message = "Username is required")  
    @Size(min = 3, max = 50)  
    private String username;  
  
    @NotBlank(message = "Email is required")  
    @Email(message = "Email should be valid")  
    private String email;  
  
    @NotBlank(message = "Password is required")  
    @Size(min = 6)  
    private String password;  
  
    @NotBlank(message = "Full name is required")  
    private String fullName;  
}
```

Authentication DTOs (cont.)

LoginRequest

```
@Data  
@Builder  
@NoArgsConstructor  
@AllArgsConstructor  
public class LoginRequest {  
    @NotBlank(message = "Username is required")  
    private String username;  
  
    @NotBlank(message = "Password is required")  
    private String password;  
}
```

Key Features:

- Bean Validation annotations
- Clear validation messages
- Type-safe data transfer

AuthResponse

```
@Data  
@Builder  
@NoArgsConstructor  
@AllArgsConstructor  
public class AuthResponse {  
    private String token;  
    private String type = "Bearer";  
    private Long userId;  
    private String username;  
    private String email;  
    private String fullName;  
}
```

Returns:

- JWT token for authentication
- User information
- Token type for HTTP header

Generic API Response

```
@Data  
@Builder  
@NoArgsConstructor  
@AllArgsConstructor  
public class ApiResponse<T> {  
    private boolean success;  
    private String message;  
    private T data;  
  
    public static <T> ApiResponse<T> success(String message, T data) {  
        return ApiResponse.<T>builder()  
            .success(true)  
            .message(message)  
            .data(data)  
            .build();  
    }  
  
    public static <T> ApiResponse<T> error(String message) {  
        return ApiResponse.<T>builder()  
            .success(false)  
            .message(message)  
            .build();  
    }  
}
```

Why Use DTOs?

Benefits

- ✓ **Separation of Concerns:** API contracts separate from entities
- ✓ **Security:** Don't expose internal entity structure
- ✓ **Validation:** Input validation at API boundary
- ✓ **Flexibility:** Change entities without breaking API
- ✓ **Performance:** Transfer only needed data

Step 4: Exception Handling

Clean Error Management

Custom Exceptions

```
// ResourceNotFoundException.java
public class ResourceNotFoundException extends RuntimeException {
    public ResourceNotFoundException(String message) {
        super(message);
    }
}

// DuplicateResourceException.java
public class DuplicateResourceException extends RuntimeException {
    public DuplicateResourceException(String message) {
        super(message);
    }
}

// InsufficientBalanceException.java
public class InsufficientBalanceException extends RuntimeException {
    public InsufficientBalanceException(String message) {
        super(message);
    }
}
```

Global Exception Handler

```
@RestControllerAdvice
public class GlobalExceptionHandler {

    @ExceptionHandler(ResourceNotFoundException.class)
    public ResponseEntity<ApiResponse<Void>> handleResourceNotFound(
        ResourceNotFoundException ex) {
        return ResponseEntity
            .status(HttpStatus.NOT_FOUND)
            .body(ApiResponse.error(ex.getMessage()));
    }

    @ExceptionHandler(DuplicateResourceException.class)
    public ResponseEntity<ApiResponse<Void>> handleDuplicateResource(
        DuplicateResourceException ex) {
        return ResponseEntity
            .status(HttpStatus.CONFLICT)
            .body(ApiResponse.error(ex.getMessage()));
    }
}
```

Validation Exception Handler

```
@ExceptionHandler(MethodArgumentNotValidException.class)
public ResponseEntity<ApiResponse<Map<String, String>>>
    handleValidation(MethodArgumentNotValidException ex) {

    Map<String, String> errors = new HashMap<>();
    ex.getBindingResult().getAllErrors().forEach((error) -> {
        String fieldName = ((FieldError) error).getField();
        String errorMessage = error.getDefaultMessage();
        errors.put(fieldName, errorMessage);
    });

    return ResponseEntity
        .status(HttpStatus.BAD_REQUEST)
        .body(ApiResponse.<Map<String, String>>builder()
            .success(false)
            .message("Validation failed")
            .data(errors)
            .build());
}
```

Benefits of `@RestControllerAdvice`

Why Use It?

- ✓ **Centralized:** All exception handling in one place
- ✓ **Clean Controllers:** No try-catch blocks needed
- ✓ **Consistent:** Same error response format everywhere
- ✓ **Maintainable:** Easy to add new exception types
- ✓ **Professional:** Proper HTTP status codes

Step 5: Repository Layer

Data Access with JPA

UserRepository

```
@Repository
public interface UserRepository extends JpaRepository<User, Long> {
    Optional<User> findByUsername(String username);
    Optional<User> findByEmail(String email);
    boolean existsByUsername(String username);
    boolean existsByEmail(String email);
}
```

Features:

- Spring Data JPA - no implementation needed!
- Type-safe queries
- Automatic CRUD operations

AccountRepository

```
@Repository
public interface AccountRepository extends JpaRepository<Account, Long> {
    Optional<Account> findByAccountNumber(String accountNumber);
    List<Account> findByUserId(Long userId);
    boolean existsByAccountNumber(String accountNumber);
}
```

Features:

- Custom query methods
- Derived queries from method names
- Spring generates implementation automatically

TransactionRepository

```
@Repository
public interface TransactionRepository
    extends JpaRepository<Transaction, Long> {

    Page<Transaction> findByAccountId(Long accountId, Pageable pageable);

    @Query("SELECT t FROM Transaction t WHERE t.account.id = :accountId " +
        "AND (CAST(:startDate AS timestamp) IS NULL " +
        "OR t.createdAt >= :startDate) " +
        "AND (CAST(:endDate AS timestamp) IS NULL " +
        "OR t.createdAt <= :endDate)")
    Page<Transaction> findByAccountIdAndDateRange(
        @Param("accountId") Long accountId,
        @Param("startDate") LocalDateTime startDate,
        @Param("endDate") LocalDateTime endDate,
        Pageable pageable
    );
}
```

Important: CAST for Nullable Parameters

Why CAST is Necessary

```
// ✗ This FAILS with PostgreSQL
"WHERE :startDate IS NULL OR t.createdAt >= :startDate"

// ✓ This WORKS
"WHERE CAST(:startDate AS timestamp) IS NULL OR t.createdAt >= :startDate"
```

Reason: PostgreSQL cannot determine the data type of nullable parameters in JPQL queries without explicit casting.

Error Without CAST: could not determine data type of parameter \$2

Step 6: JWT Authentication

Secure Token-Based Auth

JWT Utility Class

```
@Component
public class JwtUtil {
    @Value("${jwt.secret}")
    private String secret;

    @Value("${jwt.expiration}")
    private Long expiration;

    private SecretKey getSigningKey() {
        return Keys.hmacShaKeyFor(secret.getBytes(StandardCharsets.UTF_8));
    }

    @SneakyThrows
    public String generateToken(String username, Long userId) {
        Map<String, Object> claims = new HashMap<>();
        claims.put("userId", userId);
        return Jwts.builder()
            .claims(claims)
            .subject(username)
```

JWT Utility Class (cont.)

```
.issuedAt(new Date(System.currentTimeMillis()))
.expiration(new Date(System.currentTimeMillis() + expiration))
.signWith(getSigningKey())
.compact();
}

@sneakyThrows
public String extractUsername(String token) {
    return Jwts.parser()
        .verifyWith(getSigningKey())
        .build()
        .parseSignedClaims(token)
        .getPayload()
        .getSubject();
}
```

Note: `@SneakyThrows` from Lombok keeps code clean!

JWT Validation

```
@SneakyThrows
public Boolean validateToken(String token, String username) {
    final String extractedUsername = extractUsername(token);
    Date expiration = Jwts.parser()
        .verifyWith(getSigningKey())
        .build()
        .parseSignedClaims(token)
        .getPayload()
        .getExpiration();

    return (extractedUsername.equals(username) &&
        !expiration.before(new Date()));
}
```

Validates:

- Token signature
- Username matches
- Token not expired

JWT Authentication Filter

```
@Component
@RequiredArgsConstructor
public class JwtAuthenticationFilter extends OncePerRequestFilter {
    private final JwtUtil jwtUtil;

    @Override
    @SneakyThrows
    protected void doFilterInternal(HttpServletRequest request,
                                    HttpServletResponse response,
                                    FilterChain filterChain) {
        final String authHeader = request.getHeader("Authorization");

        if (authHeader == null || !authHeader.startsWith("Bearer ")) {
            filterChain.doFilter(request, response);
            return;
        }

        final String jwt = authHeader.substring(7);
        final String username = jwtUtil.extractUsername(jwt);
```

JWT Authentication Filter (cont.)

```
if (username != null &&
    SecurityContextHolder.getContext().getAuthentication() == null) {

    if (jwtUtil.validateToken(jwt, username)) {
        UsernamePasswordAuthenticationToken authToken =
            new UsernamePasswordAuthenticationToken(
                username, null, new ArrayList<>()
            );
        authToken.setDetails(
            new WebAuthenticationDetailsSource().buildDetails(request)
        );
        SecurityContextHolder.getContext()
            .setAuthentication(authToken);
    }
}
filterChain.doFilter(request, response);
}
```

Security Configuration

```
@Configuration
@EnableWebSecurity
@RequiredArgsConstructor
public class SecurityConfig {
    private final JwtAuthenticationFilter jwtAuthenticationFilter;

    @Bean
    public SecurityFilterChain securityFilterChain(HttpSecurity http)
        throws Exception {
        http
            .csrf(AbstractHttpConfigurer::disable)
            .authorizeHttpRequests(auth -> auth
                .requestMatchers("/api/auth/**").permitAll()
                .requestMatchers("/actuator/**").permitAll()
                .anyRequest().authenticated()
            )
    }
}
```

Security Configuration (cont.)

```
        .sessionManagement(session -> session
            .sessionCreationPolicy(SessionCreationPolicy.STATELESS)
        )
        .addFilterBefore(jwtAuthenticationFilter,
            UsernamePasswordAuthenticationFilter.class);
    return http.build();
}

@Bean
public PasswordEncoder passwordEncoder() {
    return new BCryptPasswordEncoder();
}
}
```

Key Points:

- Stateless sessions (JWT-based)
- Public auth endpoints
- BCrypt password hashing

Step 7: Logic Layer

Business Logic Implementation

TransactionLogic - Deposit

```
@Service
@RequiredArgsConstructor
public class TransactionLogic {
    private final TransactionRepository transactionRepository;
    private final AccountRepository accountRepository;
    private final UserRepository userRepository;

    @Transactional
    @SneakyThrows
    public TransactionResponse deposit(String username,
                                       TransactionRequest request) {
        validateAccountOwnership(username, request.getAccountNumber());

        Account account = accountRepository
            .findByAccountNumber(request.getAccountNumber())
            .orElseThrow(() ->
                new ResourceNotFoundException("Account not found"));
    }
}
```

TransactionLogic - Deposit (cont.)

```
BigDecimal balanceBefore = account.getBalance();
BigDecimal balanceAfter = balanceBefore.add(request.getAmount());

account.setBalance(balanceAfter);
accountRepository.save(account);

Transaction transaction = Transaction.builder()
    .account(account)
    .type(Transaction.TransactionType.DEPOSIT)
    .amount(request.getAmount())
    .balanceBefore(balanceBefore)
    .balanceAfter(balanceAfter)
    .description(request.getDescription())
    .referenceNumber(generateReferenceNumber())
    .build();

transaction = transactionRepository.save(transaction);
return mapToResponse(transaction);
}
```

Account Ownership Validation

```
@SneakyThrows
private void validateAccountOwnership(String username,
                                      String accountNumber) {
    User user = userRepository.findByUsername(username)
        .orElseThrow(() ->
            new ResourceNotFoundException("User not found"));

    Account account = accountRepository
        .findByAccountNumber(accountNumber)
        .orElseThrow(() ->
            new ResourceNotFoundException("Account not found"));

    if (!account.getUser().getId().equals(user.getId())) {
        throw new ResourceNotFoundException("Account not found");
    }
}
```

Security: Users can only access their own accounts!

Reference Number Generation

```
private String generateReferenceNumber() {  
    return "TXN" + UUID.randomUUID()  
        .toString()  
        .replace("-", "")  
        .substring(0, 16)  
        .toUpperCase();  
}
```

Example: TXN1A2B3C4D5E6F7G8H

Purpose: Unique identifier for audit trail

Key Logic Layer Concepts

Important Annotations

@Transactional: Ensures atomicity

- All operations succeed or all fail
- Automatic rollback on exceptions
- Database consistency guaranteed

@SneakyThrows: Clean code

- Avoids cluttering with try-catch
- Lombok annotation
- Better readability

Step 8: Route Layer

Controllers & API Endpoints

TransactionController

```
@RestController
@RequestMapping("/api/transactions")
@RequiredArgsConstructor
public class TransactionController {
    private final TransactionLogic transactionLogic;

    @PostMapping("/deposit")
    public ResponseEntity<ApiResponse<TransactionResponse>> deposit(
        Authentication authentication,
        @Valid @RequestBody TransactionRequest request) {

        TransactionResponse response = transactionLogic.deposit(
            authentication.getName(),
            request
        );

        return ResponseEntity
            .status(HttpStatus.CREATED)
            .body(ApiResponse.success("Deposit successful", response));
    }
}
```

Transaction History Endpoint

```
@GetMapping
public ResponseEntity<ApiResponse<Page<TransactionResponse>>>
    getTransactionHistory(
        Authentication authentication,
        @RequestParam String accountNumber,
        @RequestParam(required = false) LocalDateTime startDate,
        @RequestParam(required = false) LocalDateTime endDate,
        @RequestParam(defaultValue = "0") int page,
        @RequestParam(defaultValue = "10") int size,
        @RequestParam(defaultValue = "createdAt") String sortBy,
        @RequestParam(defaultValue = "desc") String sortDirection) {

    Page<TransactionResponse> response =
        transactionLogic.getTransactionHistory(
            authentication.getName(), accountNumber,
            startDate, endDate, page, size, sortBy, sortDirection
        );
}
```

Transaction History Response

```
        return ResponseEntity.ok(
            ApiResponse.success("Transaction history retrieved", response)
        );
    }
}
```

Features:

- Pagination support
- Date range filtering
- Sorting capabilities
- Authentication required

Controller Responsibilities

What Controllers Should Do

- Receive HTTP requests
- Extract authentication information
- Validate input with `@Valid`
- Call business logic layer
- Return formatted responses

What Controllers Should NOT Do

- Business logic
- Database operations
- Complex calculations

Step 9: Docker Deployment

Containerization with Docker

Why Docker?

Benefits of Containerization

- ✓ **Consistency:** Same environment everywhere
- ✓ **Portability:** Run anywhere Docker is installed
- ✓ **Isolation:** Dependencies contained
- ✓ **Scalability:** Easy to scale with orchestration
- ✓ **Simplicity:** One command to run

Dockerfile - Multi-Stage Build

Stage 1: Build

```
# Stage 1: Build the application
FROM maven:3.9.6-eclipse-temurin-17 AS build

WORKDIR /app

# Copy pom.xml and download dependencies (cached layer)
COPY pom.xml .
RUN mvn dependency:go-offline -B

# Copy source code
COPY src ./src

# Build the application (skip tests for faster builds)
RUN mvn clean package -DskipTests
```

Dockerfile - Multi-Stage Build (cont.)

Stage 2: Run

```
# Stage 2: Run the application
FROM eclipse-temurin:17-jre

WORKDIR /app

# Create a non-root user for security
RUN groupadd -r spring && useradd -r -g spring spring

# Copy the JAR from build stage
COPY --from=build /app/target/*.jar app.jar

# Change ownership to spring user
RUN chown -R spring:spring /app

# Switch to non-root user
USER spring:spring

# Expose port 8000
EXPOSE 8000
```

Dockerfile - Health Check & Entrypoint

```
# Health check (using curl)
HEALTHCHECK --interval=30s --timeout=3s --start-period=60s --retries=3 \
  CMD curl -f http://localhost:8000/actuator/health || exit 1

# JVM configuration for containerized environment
ENV JAVA_OPTS="-XX:+UseContainerSupport -XX:MaxRAMPercentage=75.0 -XX:InitialRAMPercentage=50.0"

# Run the application
ENTRYPOINT ["sh", "-c", "java $JAVA_OPTS -Djava.security.egd=file:/dev/./urandom -jar app.jar"]
```

Key Features:

- Multi-stage build (smaller image)
- Non-root user (security)
- Health check (monitoring)
- JVM optimization for containers

Docker Compose Configuration

```
services:  
  # Spring Boot Application  
  app:  
    build:  
      context: .  
      dockerfile: Dockerfile  
    container_name: bank-app  
    restart: unless-stopped  
    env_file:  
      - .env  
    environment:  
      APP_PORT: 8000  
    ports:  
      - "8000:8000"  
    networks:  
      app-net: {}  
    healthcheck:  
      test: ["CMD", "wget", "--no-verbose", "--tries=1", "--spider", "http://localhost:8000/actuator/health"]  
      interval: 30s  
      timeout: 10s  
      retries: 5  
      start_period: 60s
```

Docker Compose (cont.)

```
networks:  
  app-net:  
    external: true  
    name: "user-account-net"  
  
volumes:  
  postgres_data:  
    driver: local
```

Key Features:

- Environment variable support (`.env` file)
- Health checks
- Automatic restart
- Network isolation
- Volume persistence

Docker Commands

Build and Run

```
# Build and start the application
docker compose up --build -d

# Check container status
docker compose ps

# View logs
docker compose logs -f app

# Stop the application
docker compose down

# Remove volumes (clean start)
docker compose down -v
```

Docker Best Practices

Security

- Non-root user:** Run as `spring:spring` user
- Multi-stage build:** Smaller attack surface
- No secrets in image:** Use environment variables

Performance

- Layer caching:** Dependencies downloaded separately
- JVM optimization:** Container-aware memory settings
- Health checks:** Automatic restart on failure

Monitoring

- Actuator integration:** Health check endpoint
- Log aggregation:** `docker compose logs`

Step 10: Testing the Application

Running & Testing

Setup Database

```
-- Create database
CREATE DATABASE bank_app;

-- PostgreSQL should be running on port 6435
-- (or adjust in application.properties)
```

Environment Variables Setup

Create a `.env` file in the project root:

```
# .env file
SPRING_DATASOURCE_URL=jdbc:postgresql://localhost:6435/bank_app
SPRING_DATASOURCE_USERNAME=pg
SPRING_DATASOURCE_PASSWORD=p@ssw0rd1234
APP_PORT=8000
JWT_SECRET=your_256_byte_secret_here
JWT_EXPIRATION=86400000
```

Generate JWT Secret: `openssl rand -hex 256`

Run the Application

Option 1: Using Maven with Environment Variables

```
# Load environment variables and run
export $(cat .env | xargs) && ./mvnw spring-boot:run

# Or use the provided script
./run.sh
```

Option 2: Using Docker Compose

```
# Build and run with Docker
docker compose up --build -d

# Check status
docker compose ps

# View logs
docker compose logs -f app

# Stop the application
docker compose down
```

Application starts on: <http://localhost:8000>

Test Endpoints - Register

```
### Register a new user
POST http://localhost:8000/api/auth/register
Content-Type: application/json

{
  "username": "john_doe",
  "email": "john@example.com",
  "password": "password123",
  "fullName": "John Doe"
}
```

Response:

```
{
  "success": true,
  "message": "User registered successfully",
  "data": {
    "token": "eyJhbGciOiJIUzUxMiJ9...",
    "type": "Bearer",
    "userId": 1,
    "username": "john_doe"
  }
}
```

Test Endpoints - Login

```
### Login
POST http://localhost:8000/api/auth/login
Content-Type: application/json

{
  "username": "john_doe",
  "password": "password123"
}
```

Save the token from response for subsequent requests!

Test Endpoints - Create Account

```
### Create account
POST http://localhost:8000/api/accounts
Content-Type: application/json
Authorization: Bearer YOUR_TOKEN_HERE

{
  "accountName": "Savings Account",
  "initialBalance": 1000.00
}
```

Response:

```
{
  "success": true,
  "message": "Account created successfully",
  "data": {
    "accountNumber": "ACC1234567890",
    "accountName": "Savings Account",
    "balance": 1000.00
  }
}
```

Test Endpoints - Deposit

```
### Deposit money
POST http://localhost:8000/api/transactions/deposit
Content-Type: application/json
Authorization: Bearer YOUR_TOKEN_HERE

{
  "accountNumber": "ACC1234567890",
  "amount": 500.00,
  "description": "Salary deposit"
}
```

Test Endpoints - Transfer

```
### Transfer money
POST http://localhost:8000/api/transactions/transfer
Content-Type: application/json
Authorization: Bearer YOUR_TOKEN_HERE

{
  "fromAccountNumber": "ACC1234567890",
  "toAccountNumber": "ACC9876543210",
  "amount": 200.00,
  "description": "Payment to friend"
}
```

Test Endpoints - Transaction History

```
### Get transaction history with filters
GET http://localhost:8000/api/transactions?accountNumber=ACC1234567890&page=0&size=10&sortBy=createdAt&sortDirection=desc
Authorization: Bearer YOUR_TOKEN_HERE

### With date range filter
GET http://localhost:8000/api/transactions?accountNumber=ACC1234567890&startDate=2025-01-01T00:00:00&endDate=2025-12-31T23:59:59
Authorization: Bearer YOUR_TOKEN_HERE
```

Step 11: Spring Boot Actuator

Production-Ready Monitoring

What is Actuator?

Production-Ready Features

Spring Boot Actuator provides:

-  **Health checks**
-  **Metrics collection**
-  **Application information**
-  **Environment details**
-  **HTTP trace**
-  **Configuration properties**

Purpose: Monitor and manage your application in production

Available Actuator Endpoints

```
### Health Check
GET http://localhost:8000/actuator/health

### Application Info
GET http://localhost:8000/actuator/info

### All Available Metrics
GET http://localhost:8000/actuator/metrics

### Specific Metrics
GET http://localhost:8000/actuator/metrics/http.server.requests
GET http://localhost:8000/actuator/metrics/jvm.memory.used
GET http://localhost:8000/actuator/metrics/system.cpu.usage
```

Database Connection Metrics

```
### HikariCP Connection Pool Metrics
GET http://localhost:8000/actuator/metrics/hikaricp.connections.active
GET http://localhost:8000/actuator/metrics/hikaricp.connections.idle
GET http://localhost:8000/actuator/metrics/hikaricp.connections.max
GET http://localhost:8000/actuator/metrics/hikaricp.connections.min
```

Why Monitor Connections?

- Track connection pool usage
- Identify connection leaks
- Optimize pool configuration
- Prevent database bottlenecks

Environment & Configuration

```
### Environment Variables and Properties  
GET http://localhost:8000/actuator/env  
  
### Application Beans  
GET http://localhost:8000/actuator/beans  
  
### Request Mappings (All Endpoints)  
GET http://localhost:8000/actuator/mappings
```

Health Check Response Example

```
{  
  "status": "UP",  
  "components": {  
    "db": {  
      "status": "UP",  
      "details": {  
        "database": "PostgreSQL",  
        "validationQuery": "isValid()"  
      }  
    },  
    "diskSpace": {  
      "status": "UP",  
      "details": {  
        "total": 500068036608,  
        "free": 198648901632,  
        "threshold": 10485760  
      }  
    }  
  }  
}
```

Metrics Response Example

```
{  
  "name": "http.server.requests",  
  "measurements": [  
    {  
      "statistic": "COUNT",  
      "value": 157.0  
    },  
    {  
      "statistic": "TOTAL_TIME",  
      "value": 3.456789  
    }  
  ],  
  "availableTags": [  
    {  
      "tag": "method",  
      "values": ["GET", "POST", "PUT", "DELETE"]  
    }  
  ]  
}
```

Production Considerations

Security Best Practices

1. Secure Actuator Endpoints

```
.requestMatchers("/actuator/**").hasRole("ADMIN")
```

2. Limit Exposed Endpoints

```
management.endpoints.web.exposure.include=health,info,metrics
```

3. Hide Sensitive Information

```
management.endpoint.env.show-values=WHEN_AUTHORIZED
```

Production Considerations (cont.)

Integration with Monitoring Tools

Prometheus Integration:

```
<dependency>
  <groupId>io.micrometer</groupId>
  <artifactId>micrometer-registry-prometheus</artifactId>
</dependency>
```

Grafana Dashboard: Visualize metrics

ELK Stack: Log aggregation and analysis

CloudWatch: AWS monitoring

Step 12: Comprehensive Logging

Request/Response Logging with JSON Support

What is Logging?

Why Comprehensive Logging Matters

In production environments, you need to know:

- 🔎 **What happened?** - Every request and response
- 👤 **Who did it?** - User identification
- ⏳ **How long?** - Performance tracking
- ✗ **What went wrong?** - Error details with context

Our logging system provides all of this automatically!

Logging Features

Automatic Request/Response Logging

- ✓ **Timestamp:** Millisecond precision (yyyy-MM-dd HH:mm:ss.SSS)
- ✓ **User ID:** Extracted from JWT authentication
- ✓ **HTTP Method & URI:** Full request path with query params
- ✓ **Status Code:** HTTP response status
- ✓ **Elapsed Time:** Request processing time in milliseconds
- ✓ **Request/Response Body:** For errors only (with sensitive data masked)

Dual Format Support

Two Logging Formats for Different Needs

Human-Readable (Development):

```
2025-11-18 21:30:45.123 INFO - Timestamp: 2025-11-18 21:30:45.123 | User ID: john@example.com | Method: GET | URI: /api/accounts | Status: 200 | Elapsed Time: 45 ms
```

JSON Format (Production):

```
{
  "timestamp": "2025-11-18T14:30:45.123Z",
  "level": "INFO",
  "userId": "john@example.com",
  "httpMethod": "GET",
  "httpUri": "/api/accounts",
  "httpStatusCode": "200",
  "elapsedTime": "45"
}
```

Switching Between Formats

Easy Profile-Based Configuration

Development (human-readable):

```
./mvnw spring-boot:run
```

Production (JSON):

```
./mvnw spring-boot:run -Dspring-boot.run.profiles=json
```

Docker:

```
docker run -e SPRING_PROFILES_ACTIVE=json bank-app
```

Error Logging Example

Detailed Error Information

```
===== ERROR REQUEST/RESPONSE =====
Timestamp: 2025-11-18 21:31:15.456
User ID: john@example.com
Method: POST
URI: /api/accounts/123/withdraw
Status Code: 400
Elapsed Time: 12 ms
--- Request Body ---
{"amount":5000.00}
--- Response Body ---
{"success":false,"message":"Insufficient balance","data":null}
=====
```

Note: Only errors show request/response bodies. Success requests show summary only.

Security Features

Sensitive Data Masking

Automatically masks sensitive information:

Before Masking:

```
{"email":"user@example.com", "password":"mySecret123"}
```

After Masking (in logs):

```
{"email":"user@example.com", "password": "***MASKED***"}
```

Fields Masked: password , token , secret

Implementation Components

RequestResponseLoggingFilter

```
@Slf4j
@Component
@Order(Ordered.HIGHEST_PRECEDENCE)
public class RequestResponseLoggingFilter extends OncePerRequestFilter {

    @Override
    protected void doFilterInternal(HttpServletRequest request,
                                    HttpServletResponse response,
                                    FilterChain filterChain) {
        long startTime = System.currentTimeMillis();

        ContentCachingRequestWrapper wrappedRequest =
            new ContentCachingRequestWrapper(request);
        ContentCachingResponseWrapper wrappedResponse =
            new ContentCachingResponseWrapper(response);

        try {
            filterChain.doFilter(wrappedRequest, wrappedResponse);
        } finally {
            long elapsedTime = System.currentTimeMillis() - startTime;
            logRequestResponse(wrappedRequest, wrappedResponse, elapsedTime);
            wrappedResponse.copyBodyToResponse();
        }
    }
}
```

Logback Configuration

logback-spring.xml

```
<configuration>
    <!-- Human-readable appender -->
    <appender name="CONSOLE" class="ch.qos.logback.core.ConsoleAppender">
        <encoder>
            <pattern>%d{yyyy-MM-dd HH:mm:ss.SSS} %highlight(%-5level)
                [%thread] %cyan(%logger{36}) - %msg%n</pattern>
        </encoder>
    </appender>

    <!-- JSON appender -->
    <appender name="CONSOLE_JSON"
              class="ch.qos.logback.core.ConsoleAppender">
        <encoder class="net.logstash.logback.encoder.LogstashEncoder">
            <includeMdcKeyName>userId</includeMdcKeyName>
            <includeMdcKeyName>httpMethod</includeMdcKeyName>
            <includeMdcKeyName>httpUri</includeMdcKeyName>
            <includeMdcKeyName>httpStatusCode</includeMdcKeyName>
            <includeMdcKeyName>elapsedTime</includeMdcKeyName>
        </encoder>
    </appender>
```

Logback Configuration (cont.)

```
<!-- Profile: default (human-readable) -->
<springProfile name="default,dev,local">
    <root level="INFO">
        <appender-ref ref="CONSOLE"/>
    </root>
</springProfile>

<!-- Profile: json (JSON for production) -->
<springProfile name="json,prod,production">
    <root level="INFO">
        <appender-ref ref="CONSOLE_JSON"/>
    </root>
</springProfile>
</configuration>
```

Key: Use Spring profiles to switch between formats!

Log Aggregation Integration

Ready for Production Monitoring

ELK Stack (Elasticsearch, Logstash, Kibana):

```
input {  
    file {  
        path => "/app/logs/application.json"  
        codec => "json"  
    }  
}  
output {  
    elasticsearch {  
        hosts => ["localhost:9200"]  
        index => "bank-app-%{+YYYY.MM.dd}"  
    }  
}
```

Also Compatible With:

- AWS CloudWatch
- Splunk
- Datadog
- Grafana Loki

Querying JSON Logs

Using jq for Analysis

```
# Get all error logs
cat logs/application.json | jq 'select(.level == "ERROR")'

# Get logs for specific user
cat logs/application.json | jq 'select(.userId == "john@example.com")'

# Get slow requests (> 1000ms)
cat logs/application.json | jq 'select(.elapsedTime | tonumber > 1000)'

# Count requests per endpoint
cat logs/application.json | jq -r '.httpUri' | sort | uniq -c

# Calculate average response time
cat logs/application.json | jq -r '.elapsedTime' | \
awk '{sum+=$1; n++} END {print sum/n}'
```

Logging Benefits

Why This Matters in Production

- ✓ **Debugging:** Find exactly what went wrong and why
- ✓ **Performance Monitoring:** Identify slow endpoints
- ✓ **Security Auditing:** Track who accessed what
- ✓ **Business Analytics:** API usage patterns
- ✓ **Compliance:** Audit trail for regulations
- ✓ **Alerting:** Trigger alerts on errors or slow requests

Documentation

Comprehensive Guides Available

-  **LOGGING_README.md:** Quick start guide
-  **LOGGING_GUIDE.md:** Human-readable logging
-  **JSON_LOGGING_GUIDE.md:** JSON logging & integration
-  **TEST_LOGGING.md:** Testing the logging system
-  **LOGGING_EXAMPLES.md:** Format comparisons

Key Features Implemented

What We've Built

Feature 1: Clean Architecture

Three-Layered Design

- Route Layer:** HTTP handling & validation
- Logic Layer:** Business rules & transactions
- Service Layer:** Data access & repositories

Benefits:

- Clear separation of concerns
- Easy to test independently
- Maintainable and scalable
- Professional code organization

Feature 2: JWT Authentication

Secure Token-Based Auth

- Stateless sessions (no server-side storage)
- Secure token generation with HMAC-SHA
- Token validation on every request
- Protected endpoints with Spring Security

Benefits:

- Scalable (no session storage)
- Secure (encrypted tokens)
- RESTful (stateless)
- Industry standard

Feature 3: Global Exception Handling

@RestControllerAdvice

- Centralized error management
- Consistent error responses
- Proper HTTP status codes
- Validation error handling

Benefits:

- Clean controller code
- Professional API responses
- Easy to maintain
- Better debugging

Feature 4: SneakyThrows Usage

Lombok's @SneakyThrows

```
// Without @SneakyThrows
public void deposit() {
    try {
        // logic
    } catch (Exception e) {
        throw new RuntimeException(e);
    }
}

// With @SneakyThrows
@sneakyThrows
public void deposit() {
    // logic
}
```

Benefits: Cleaner code, better readability

Feature 5: Transaction Management

@Transactional

- Atomicity:** All or nothing
- Consistency:** Data integrity maintained
- Isolation:** Concurrent transaction handling
- Durability:** Permanent once committed

Example: Transfer operation - both debit and credit succeed, or both fail!

Feature 6: Pagination & Filtering

Efficient Data Retrieval

- Page-based results (avoid loading all data)
- Custom page size
- Date range filtering
- Sorting capabilities (ASC/DESC)

Example:

```
GET /api/transactions?page=0&size=10&sortBy=createdAt&sortDirection=desc
```

Feature 7: Spring Boot Actuator

Production Monitoring

- Health checks
- Performance metrics
- Database connection monitoring
- Application insights

Benefits: Know what's happening in production!

Feature 8: HikariCP Connection Pooling

High-Performance Database Connections

Connection pooling (reuse connections)

Configurable pool size

Connection timeout handling

Leak detection

Configuration:

- Min idle: 5
- Max pool size: 20
- Connection timeout: 30s
- Idle timeout: 5m

Feature 8: Comprehensive Logging System

Production-Ready Logging

- Dual format support (human-readable + JSON)
- Automatic request/response logging
- Sensitive data masking
- Performance tracking (elapsed time)
- User identification from JWT
- Error details with request/response bodies

Benefits: Debug issues, monitor performance, audit access, integrate with log aggregation systems

Best Practices Demonstrated

Professional Spring Boot Development

Best Practice 1: DTOs for Data Transfer

Separate Entities from API Contracts

- Security:** Don't expose internal structure
- Flexibility:** Change entities without breaking API
- Validation:** Validate at API boundary
- Documentation:** Clear API contracts

Example: UserEntity has password, but UserResponse does not!

Best Practice 2: Bean Validation

Automatic Input Validation

```
@NotBlank(message = "Username is required")
@Size(min = 3, max = 50)
private String username;

@email(message = "Email should be valid")
private String email;
```

- Declarative validation
- Consistent error messages
- No boilerplate validation code

Best Practice 3: Builder Pattern

Clean Object Creation with Lombok

```
User user = User.builder()  
    .username("john_doe")  
    .email("john@example.com")  
    .fullName("John Doe")  
    .build();
```

- Readable code
- Immutable objects
- Optional parameters
- No constructor pollution

Best Practice 4: Repository Pattern

Abstract Data Access

```
public interface UserRepository extends JpaRepository<User, Long> {  
    Optional<User> findByUsername(String username);  
}
```

- No SQL code needed
- Type-safe queries
- Easy to test (mockable)
- Switch databases easily

Best Practice 5: Dependency Injection

Loose Coupling with Constructor Injection

```
@RequiredArgsConstructor  
public class TransactionLogic {  
    private final TransactionRepository transactionRepository;  
    private final AccountRepository accountRepository;  
}
```

- Immutable** (final fields)
- Testable** (inject mocks)
- Explicit dependencies**
- Thread-safe**

Best Practice 6: RESTful Design

Standard HTTP Methods & Status Codes

Operation	Method	Status Code
Create	POST	201 Created
Read	GET	200 OK
Update	PUT/PATCH	200 OK
Delete	DELETE	204 No Content
Error	Any	400/404/500

 Predictable API

 Industry standard

 Easy to understand

Best Practice 7: Security

Multiple Security Layers

- JWT:** Stateless authentication
- BCrypt:** Password hashing
- Ownership Validation:** Users access only their data
- Input Validation:** Prevent injection attacks
- HTTPS:** Encrypt data in transit (production)

Project Structure Recap

Clean Organization

Complete Project Structure

```
src/main/java/com/user/account/app/
  config/
    JwtUtil.java
    JwtAuthenticationFilter.java
    SecurityConfig.java
    RequestResponseLoggingFilter.java
    CachedBodyHttpServletRequest.java
    CachedBodyHttpServletResponse.java
  dto/
    RegisterRequest.java
    LoginRequest.java
    AuthResponse.java
    AccountRequest.java
    AccountResponse.java
    TransactionRequest.java
    TransactionResponse.java
    ApiResponse.java
```

Project Structure (cont.)

```
entities/
├── User.java
├── Account.java
└── Transaction.java
exceptions/
├── ResourceNotFoundException.java
├── DuplicateResourceException.java
├── InsufficientBalanceException.java
└── GlobalExceptionHandler.java
logics/
├── AuthLogic.java
├── AccountLogic.java
└── TransactionLogic.java
```

Project Structure (cont.)

```
routes/
├── AuthController.java
├── AccountController.java
└── TransactionController.java
services/
├── UserRepository.java
├── AccountRepository.java
└── TransactionRepository.java
UserAccountApplication.java
```

Conclusion

What We've Accomplished

Summary: What We Built

A production-ready banking application with:

- ✓ Clean three-layered architecture
- ✓ Secure JWT authentication
- ✓ Comprehensive exception handling
- ✓ Transaction management with ACID guarantees
- ✓ Pagination and filtering capabilities
- ✓ Spring Boot Actuator monitoring
- ✓ HikariCP connection pooling
- ✓ PostgreSQL with optimized queries
- ✓ RESTful API design
- ✓ Docker containerization with multi-stage builds
- ✓ Environment variable configuration
- ✓ **Comprehensive logging system (human-readable + JSON)**
- ✓ **Log aggregation ready (ELK, Splunk, CloudWatch)**

Why This Architecture?

The Benefits

Maintainability: Easy to locate and modify code

Testability: Each layer tested independently

Scalability: Easy to add new features

Readability: Clear organization and purpose

Professional: Industry best practices

Next Steps

Extending the Application

Suggested Enhancements

Feature Additions

1. **User Roles & Permissions:** Admin, Manager, Customer
2. **Account Statements:** PDF generation
3. **Email Notifications:** Transaction alerts
4. **Transaction Limits:** Daily/monthly limits
5. **Audit Logging**  Implemented with logging system!
6. **API Rate Limiting:** Prevent abuse
7. **Swagger/OpenAPI:** API documentation

Additional Improvements

Technical Enhancements

1. **Redis Caching:** Improve performance
2. **Message Queue:** Async processing (RabbitMQ/Kafka)
3. ~~Docker~~: Containerization  Implemented!
4. **Kubernetes:** Orchestration
5. **CI/CD Pipeline:** Automated deployment
6. **Integration Tests:** Full API testing
7. **Load Testing:** Performance verification

Resources

Documentation

Project Files

-  **API_DOCUMENTATION.md:** Detailed API specifications
-  **api-tests.http:** Testing examples
-  **MEDIUM_ARTICLE.md:** This tutorial
-  **application.properties:** Configuration reference
-  **LOGGING_README.md:** Logging quick start
-  **JSON_LOGGING_GUIDE.md:** JSON logging integration

Testing the Application

Quick Start Commands

Option 1: Maven

```
# Build  
./mvnw clean install  
  
# Run with environment variables  
export $(cat .env | xargs) && ./mvnw spring-boot:run  
  
# Or use the script  
./run.sh
```

Option 2: Docker

```
# Build and run  
docker compose up --build -d  
  
# View logs  
docker compose logs -f app
```

Test the application:

```
curl http://localhost:8000/actuator/health
```

Bonus: N+1 Query Problem

Common Performance Issue & Solutions

What is the N+1 Query Problem?

The Problem

The N+1 query problem is a common performance issue in JPA/Hibernate where:

- **1 query** fetches N parent entities
- **N queries** fetch related child entities (one query per parent)

This results in **1 + N = N+1 queries** instead of a single optimized query.

Example Entities

- `Author` (parent) - has many books
- `Book` (child) - belongs to one author

The Problem in Action

```
List<Author> authors = authorRepository.findAll(); // 1 query
for (Author author : authors) {
    author.getBooks().size(); // N queries (one per author!)
}
```

SQL Executed:

```
-- Query 1: Fetch all authors
SELECT * FROM authors;

-- Query 2: For author 1
SELECT * FROM books WHERE author_id = 1;

-- Query 3: For author 2
SELECT * FROM books WHERE author_id = 2;

-- ... and so on for each author
```

With 5 authors: $1 + 5 = 6$ queries

With 100 authors: $1 + 100 = 101$ queries!

How to Identify N+1 Problems

Enable SQL Logging

```
# application.properties  
spring.jpa.show-sql=true  
spring.jpa.properties.hibernate.format_sql=true  
spring.jpa.properties.hibernate.generate_statistics=true
```

Watch Console Output

You'll see multiple SELECT queries executing:

```
Hibernate: select a1_0.id,a1_0.email,a1_0.name from authors a1_0  
Hibernate: select b1_0.author_id,b1_0.id,b1_0.isbn,b1_0.title  
           from books b1_0 where b1_0.author_id=?  
Hibernate: select b1_0.author_id,b1_0.id,b1_0.isbn,b1_0.title  
           from books b1_0 where b1_0.author_id=?  
...
```

Solution 1: FETCH JOIN (Recommended)

Use JPQL with JOIN FETCH

```
@Query("SELECT DISTINCT a FROM Author a LEFT JOIN FETCH a.books")
List<Author> findAllWithBooks();
```

SQL Executed:

```
SELECT DISTINCT a.*, b.*
FROM authors a
LEFT JOIN books b ON a.id = b.author_id;
```

Result: Only 1 query! 

When to use:

- When you ALWAYS need the associated data
- For read-heavy operations
- When the association size is reasonable

Solution 2: @EntityGraph

Dynamic Fetching Strategy

```
@Query("SELECT a FROM Author a")
@EntityGraph(attributePaths = {"books"})
List<Author> findAllWithBooksEntityGraph();
```

Result: Only 1 query! 

When to use:

- More flexible than FETCH JOIN
- Can specify multiple attribute paths
- Good for dynamic fetching strategies

Example:

```
@EntityGraph(attributePaths = {"books", "publisher", "reviews"})
List<Author> findAllWithDetails();
```

Solution 3: Batch Fetching

Configure Batch Size

```
@Entity  
 @BatchSize(size = 10)  
 public class Author {  
     @OneToMany(mappedBy = "author")  
     @BatchSize(size = 10)  
     private List<Book> books;  
 }
```

Result: Fewer queries ($1 + N/10$)

Example:

- 100 authors without batch: 101 queries
- 100 authors with batch size 10: 11 queries

When to use:

- When you can't use FETCH JOIN
- Reduces queries but doesn't eliminate them
- Good middle ground

Solution 4: DTO Projection

Only Fetch What You Need

```
@Query("SELECT new AuthorDTO(a.id, a.name, COUNT(b)) " +  
       "FROM Author a LEFT JOIN a.books b GROUP BY a.id, a.name")  
List<AuthorDTO> findAllAuthorsWithBookCount();
```

When to use:

- When you don't need full entities
- For reporting/dashboards
- Best performance

Benefits:

- Only retrieves required fields
- No lazy loading issues
- Perfect for read-only views

Performance Impact Comparison

With 100 Authors

Approach	Database Round Trips	Performance
N+1 Problem	101 queries	✗ Slowest
Batch Fetch (size=10)	11 queries	⚠ Better
FETCH JOIN	1 query	✓ Best
EntityGraph	1 query	✓ Best
DTO Projection	1 query	✓ Best

Performance improvement: ~100x faster with proper solution!

Best Practices

DO 

1. Always use **FETCH JOIN** or **EntityGraph** when you know you'll need associated data
2. Keep **lazy fetching as default** - only fetch what you need
3. Monitor **SQL logs** in development to catch N+1 issues early
4. Use **batch fetching** when **FETCH JOIN** isn't possible
5. Consider **DTOs** for read-only operations
6. Test with **realistic data volumes** - N+1 problems get worse with more data

Common Mistakes to Avoid

DON'T

1. Changing FetchType.LAZY to FetchType.EAGER globally

- This causes other performance issues
- Use FETCH JOIN instead

2. Ignoring N+1 in development

- "It works with 5 records" doesn't mean it works with 5000

3. Using findAll() when you need associations

- Always use custom queries with FETCH JOIN

4. Forgetting DISTINCT with FETCH JOIN

- Can cause duplicate results with OneToMany

Console Output Comparison

N+1 Problem Output ✗

```
Hibernate: select a1_0.id,a1_0.email,a1_0.name from authors a1_0
Hibernate: select b1_0.author_id,b1_0.id,b1_0.isbn,b1_0.title
           from books b1_0 where b1_0.author_id=?
Hibernate: select b1_0.author_id,b1_0.id,b1_0.isbn,b1_0.title
           from books b1_0 where b1_0.author_id=?
...
...
```

FETCH JOIN Solution Output ✓

```
Hibernate: select distinct a1_0.id,b1_0.author_id,b1_0.id,
           b1_0.isbn,b1_0.title,a1_0.email,a1_0.name
           from authors a1_0
           left join books b1_0 on a1_0.id=b1_0.author_id
```

Real-World Application

In Our Banking Application

```
// ✗ BAD: N+1 Problem
@Query("SELECT u FROM User u")
List<User> findAll();

// When accessing accounts: N queries
for (User user : users) {
    user.getAccounts().size(); // Triggers query per user!
}

// ✓ GOOD: FETCH JOIN
@Query("SELECT DISTINCT u FROM User u LEFT JOIN FETCH u.accounts")
List<User> findAllWithAccounts();

// Only 1 query for everything!
```

Bonus: Testing N+1 Detection

Hibernate Statistics

```
# Enable statistics
spring.jpa.properties.hibernate.generate_statistics=true
```

```
@Autowired
private SessionFactory sessionFactory;

@Test
public void testNoNPlusOne() {
    Statistics stats = sessionFactory.getStatistics();
    stats.clear();

    List<Author> authors = authorRepository.findAllWithBooks();

    // Assert only 1 query was executed
    assertEquals(1, stats.getPrepareStatementCount());
}
```

Key Takeaways: N+1 Query Problem

Remember

- 🎯 Always monitor SQL logs in development
- 🎯 Use **FETCH JOIN** or **EntityGraph** when loading associations
- 🎯 Lazy loading is good - but fetch eagerly when needed
- 🎯 Test with realistic data - problems scale with data size
- 🎯 DTOs for reports - best performance for read-only data

The Golden Rule

If you're loading a collection in a loop, you probably have an N+1 problem!

Questions?

Thank You!

Happy Coding! 

Remember:

- Clean architecture matters
- Security is not optional
- Test your code
- Monitor in production
- Keep learning!

Contact & Resources

Questions?

Check the comprehensive documentation in the repository

Code Examples

All code available at: `/src/main/java/com/user/account/app/`

Testing

Use `api-tests.http` for quick testing

Full Tutorial

Read `MEDIUM_ARTICLE.md` for detailed explanations

Now go build amazing Spring Boot applications! 